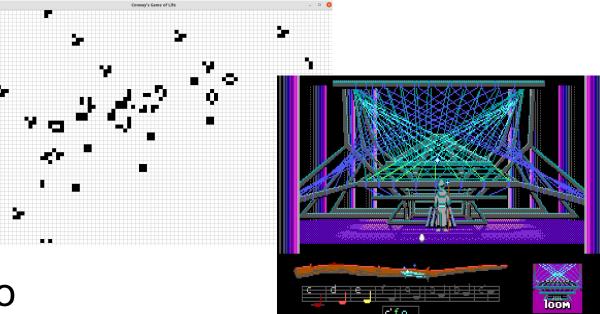


Game of Loom 2: life and dead(lock) of a virtual thread



by Mario Fusco



@mariofusco

Where did we leave? A recap on virtual threads ...

- Virtual threads make a better use of OS threads getting rid of the
 1:1 relationship with operating system threads
 - Many virtual threads can be multiplexed on the same platform thread (carrier)



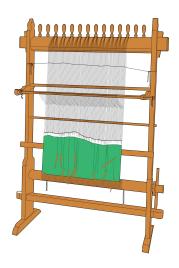
Where did we leave? A recap on virtual threads ...

- Virtual threads make a better use of OS threads getting rid of the
 1:1 relationship with operating system threads
 - Many virtual threads can be multiplexed on the same platform thread (carrier)
- When a virtual thread calls a blocking operation the JDK performs a nonblocking OS call and automatically suspends the virtual thread until the operation finishes
 - Virtual threads perform blocking calls without consuming resources, thus allowing to write asynchronous code in a way that looks synchronous
 - What virtual threads are really good for is ... waiting

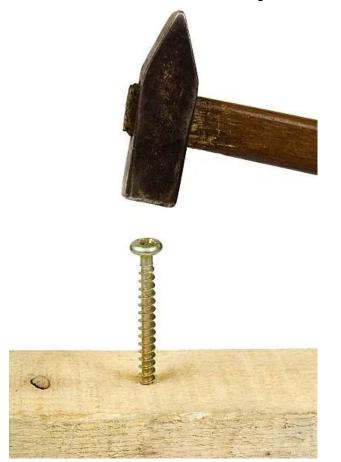


Where did we leave? A recap on virtual threads ...

- Virtual threads make a better use of OS threads getting rid of the
 1:1 relationship with operating system threads
 - Many virtual threads can be multiplexed on the same platform thread (carrier)
- When a virtual thread calls a blocking operation the JDK performs a nonblocking OS call and automatically suspends the virtual thread until the operation finishes
 - Virtual threads perform blocking calls without consuming resources, thus allowing to write asynchronous code in a way that looks synchronous
 - What virtual threads are really good for is ... waiting
- Virtual threads are lightweight
 - > 200-300B metadata
 - Pay-as-you-go stack (allocated on heap)
 - > Some ns (or below 1µs) for **context switch** (in **user space**)









Not ideal for CPU-bound task



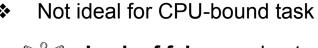
Lack of fairness due to the fact that virtual threads are not pre-emptive



Overall performance concerns caused by overhead of virtual threads scheduler









Lack of fairness due to the fact that virtual threads are not pre-emptive



Overall performance concerns caused by overhead of virtual threads scheduler

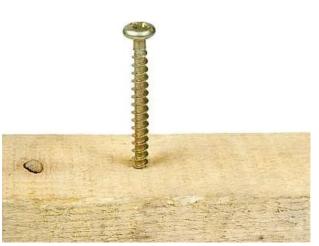
Not virtual-threads friendly frameworks & tools



Use of synchronize can lead to pinning



Use of ThreadLocal









Lack of fairness due to the fact that virtual threads are not pre-emptive



Overall performance concerns caused by overhead of virtual threads scheduler





Use of synchronize can lead to pinning

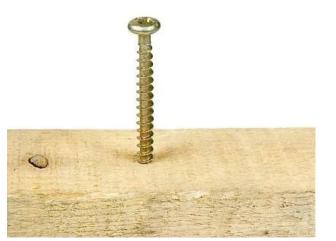


Use of ThreadLocal

Poorly customizable

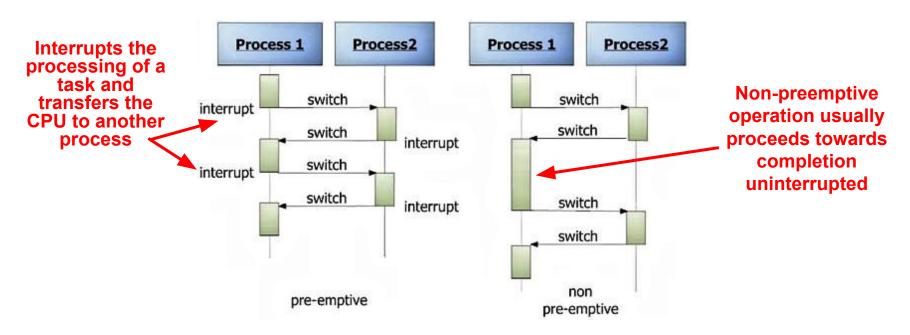


Carrier native thread pool is **not pluggable** (but for a very good reason)



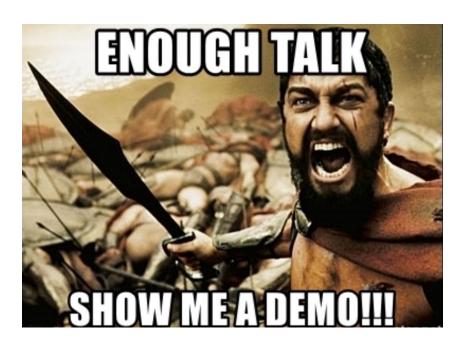
Lack of fairness in CPU bound tasks

Virtual Threads are **not pre-emptive**, they cannot be descheduled while running heavy calculations without ever calling a JDK's blocking methods, so they are **not a good fit for CPU-bound** tasks when **fairness** is key



Lack of fairness in CPU bound tasks

Virtual Threads are **not pre-emptive**, they cannot be descheduled while running heavy calculations without ever calling a JDK's blocking methods, so they are **not a good fit for CPU-bound** tasks when **fairness** is key



The cost of a continuation

In Project Loom, the word **continuation** means **delimited continuation**, also sometimes called a **coroutine**. It can be thought of as sequential code that may suspend or yield execution at some point by itself and can be resumed by a caller.

When a virtual thread hits a blocking call its continuation is **suspended** and the JVM **unmounts** the thread from its carrier ...

The cost of a continuation

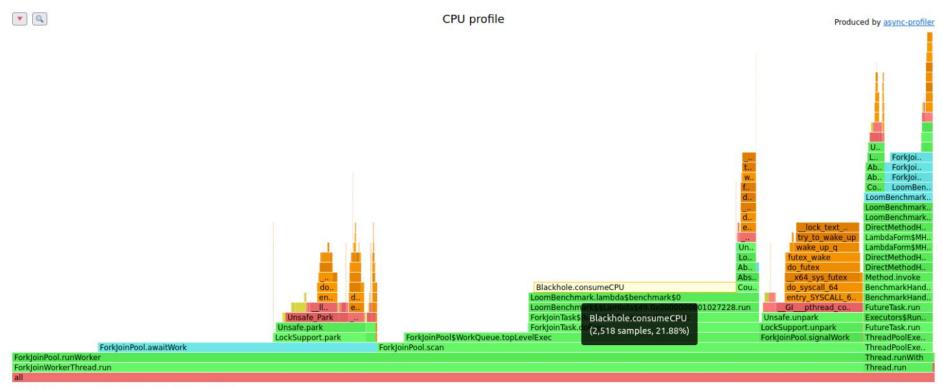
In Project Loom, the word **continuation** means **delimited continuation**, also sometimes called a **coroutine**. It can be thought of as sequential code that may suspend or yield execution at some point by itself and can be resumed by a caller.

When a virtual thread hits a blocking call its continuation is **suspended** and the JVM **unmounts** the thread from its carrier ...

... but this also has a cost!

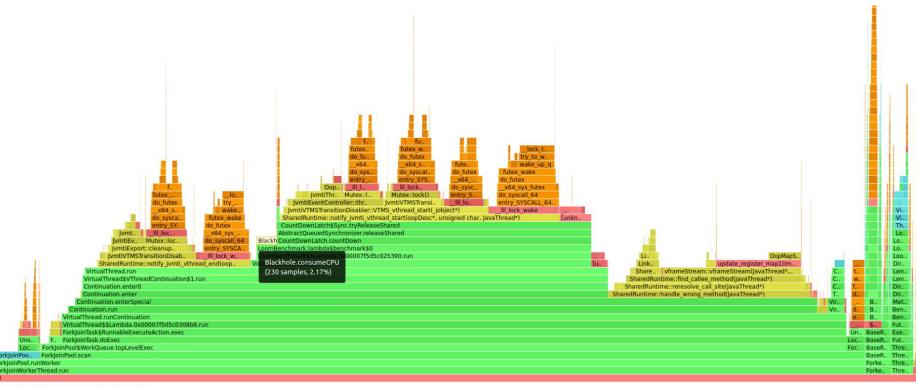


Native Threads Execution

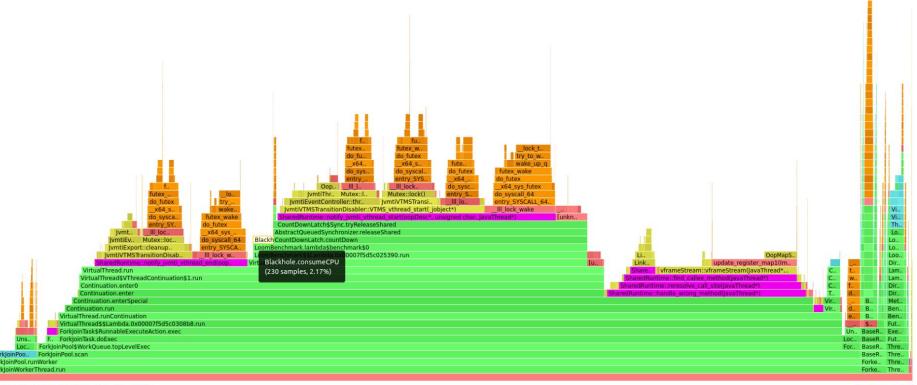


Function: Blackhole.consumeCPU (2,518 samples, 21.88%)

Virtual Threads Execution - The cost of a continuation



Virtual Threads Execution - The cost of a continuation



Function: Blackhole.consumeCPU (230 samples, 2.17%)

Matched: 70.53%

Matched: 70.

ThreadLocals as thread scoped variables

```
enum RightsLevel { ADMIN, GUEST }
record Principal(RightsLevel rights) {
   public boolean canOpen() {
       return rights == RightsLevel.ADMIN;
record Request(boolean authorized) { }
class Server {
   final static ThreadLocal<Principal> PRINCIPAL =
                              new ThreadLocal<>();
   String serve(Request request) {
       var level = request.authorized() ?
            RightsLevel.ADMIN : RightsLevel.GUEST;
       PRINCIPAL.set( new Principal(level) );
       return Application.handle( request );
```

ThreadLocals as thread scoped variables

```
enum RightsLevel { ADMIN, GUEST }
record Principal(RightsLevel rights) {
  public boolean canOpen() {
      return rights == RightsLevel.ADMIN;
record Request(boolean authorized) { }
class Server {
   final static ThreadLocal<Principal> PRINCIPAL =
                              new ThreadLocal<>();
  String serve(Request request) {
      var level = request.authorized() ?
            RightsLevel.ADMIN : RightsLevel.GUEST;
      PRINCIPAL.set( new Principal(level) );
      return Application.handle( request );
```

```
static class DBConnection {
   static DBConnection open() {
       var principal = Server.PRINCIPAL.get();
       if (!principal.canOpen())
            throw new IllegalArgumentException();
       return new DBConnection();
   String doQuery() {
       return "Result";
static class Application {
   public static String handle(Request request) {
       return DBConnection.open().doQuery();
```

ThreadLocals as thread scoped variables

```
enum RightsLevel { ADMIN, GUEST }
record Principal(RightsLevel rights) {
  public boolean canOpen() {
      return rights == RightsLevel.ADMIN;
record Request(boolean authorized) { }
class Server {
   final static ThreadLocal<Principal> PRINCIPAL =
                              new ThreadLocal<>();
  String serve(Request request) {
      var level = request.authorized() ?
            RightsLevel.ADMIN : RightsLevel.GUEST;
      PRINCIPAL.set( new Principal(level) );
      return Application.handle( request );
```

A ThreadLocal is a construct that allows us to store data accessible only by a specific thread.

```
static class DBConnection {
   static DBConnection open() {
       var principal = Server.PRINCIPAL.get();
       if (!principal.canOpen())
            throw new IllegalArgumentException();
       return new DBConnection();
   String doQuery() {
       return "Result";
static class Application {
   public static String handle(Request request) {
       return DBConnection.open().doQuery();
```

```
public static void main(String[] args) {
   var server = new Server();
   var authorized = Thread.ofVirtual().name("Authorized").unstarted(() -> callServer(server, true));
   var notAuthorized = Thread.ofVirtual().name("NOT Authorized").unstarted(() -> callServer(server, false));
   authorized.start();
   notAuthorized.start();
   try { Thread.sleep(1000L); } catch (InterruptedException e) { throw new RuntimeException(e); }
private static void callServer(Server server, boolean auth) {
   var result = server.serve(new Request(auth));
   System.out.println( "thread " + Thread.currentThread().getName() + " got result " + result );
```

```
public static void main(String[] args) {
   var server = new Server();
   var authorized = Thread.ofVirtual().name("Authorized").unstarted(() -> callServer(server, true));
   var notAuthorized = Thread.ofVirtual().name("NOT Authorized").unstarted(() -> callServer(server, false));
   authorized.start();
   notAuthorized.start();
   try { Thread.sleep(1000L); } catch (InterruptedException e) { throw new RuntimeException(e); }
private static void callServer(Server server, boolean auth) {
   var result = server.serve(new Request(auth));
   System.out.println( "thread " + Thread.currentThread().getName() + " got result " + result );
```

```
Exception in thread "NOT Authorized" java.lang.IllegalArgumentException at org.mfusco.loom.experiments.threadlocal.ThreadLocalMain$DBConnection.open(ThreadLocalMain.java:48) at org.mfusco.loom.experiments.threadlocal.ThreadLocalMain$Application.handle(ThreadLocalMain.java:41) thread Authorized got result Result
```

So? Everything is fine, right?
What is the problem to use
ThreadLocal with virtual threads?



So? Everything is fine, right?
What is the problem to use
ThreadLocal with virtual threads?





- That we can have a huge number of virtual threads, and each virtual thread will have its own ThreadLocal
 - The memory footprint of the application may quickly become very high
 - ThreadLocal will be useless in a one-thread-per-request scenario since data won't be shared between different requests

ScopedValues to the rescue

```
class Server {
    final static ThreadLocal<Principal> PRINCIPAL = new ThreadLocal<>();

String serve(Request request) {
    var level = request.authorized() ? RightsLevel.ADMIN : RightsLevel.GUEST;
    PRINCIPAL.set( new Principal(level) );
    return Application.handle( request );
  }
}
```

ScopedValues to the rescue (still preview in JDK 21)

```
class Server {
    final static ThreadLocal<Principal> PRINCIPAL = new ThreadLocal<>();

String serve(Request request) {
    var level = request.authorized() ? RightsLevel.ADMIN
    PRINCIPAL.set( new Principal(level) );
    return Application.handle( request );
  }
}
```







- It's very common to use a ThreadLocal as a thread-safe object pool
 - Easy to implement and use
 - Does not require an explicit lifecycle: it is not necessary to put the pooled resource back into the pool when finished



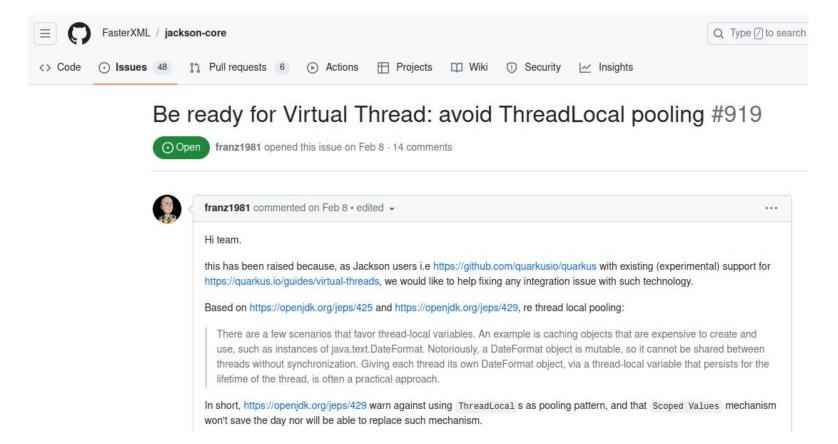
- It's very common to use a ThreadLocal as a thread-safe object pool
 - Easy to implement and use
 - Does not require an explicit lifecycle: it is not necessary to put the pooled resource back into the pool when finished
 - ScopedValues don't help in this scenario
 - This pattern plays particularly bad with virtual threads



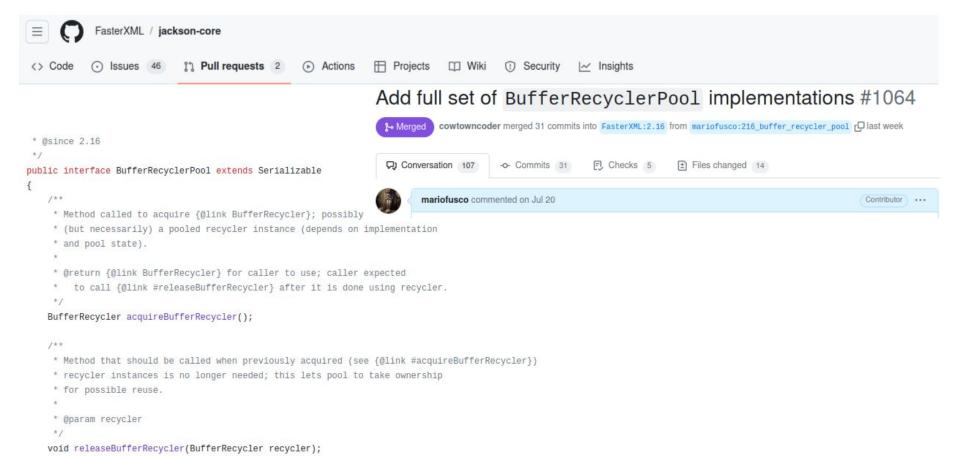
- It's very common to use a ThreadLocal as a thread-safe object pool
 - Easy to implement and use
 - Doesn't require an explicit lifecycle: it is not necessary to put the pooled resource back into the pool when finished
 - ScopedValues don't help in this scenario
 - This pattern plays particularly bad with virtual threads



ThreadLocals as Cache / Objects Pool into the wild ...



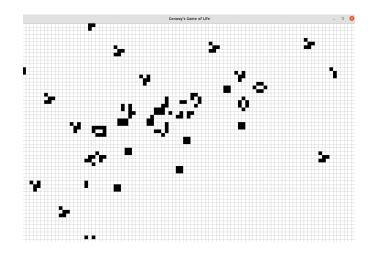
... and the fix 🎉



Can we outperform the default Fork/Join pool based carrier thread pool?

Can we outperform the default Fork/Join pool based carrier thread pool?

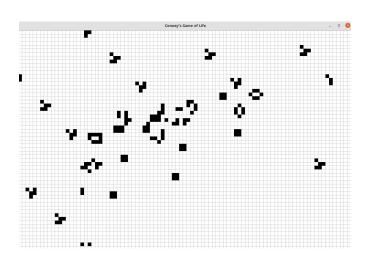
Let's put at work Game of Life again!



Can we outperform the default Fork/Join pool based carrier thread pool?

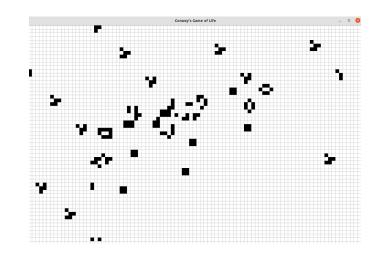
Let's put at work Game of Life again!





Can we outperform the default Fork/Join pool based carrier thread pool?

Let's put at work Game of Life again!



Error

Score

40.317 ± 0.353

244.639 ± 1.419

129.525 ± 5.346

59.071 ± 1.532

Units

ops/s

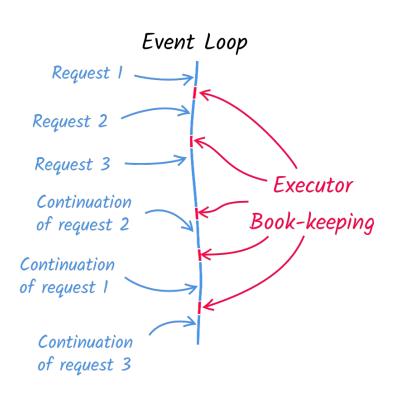
ops/s

ops/s

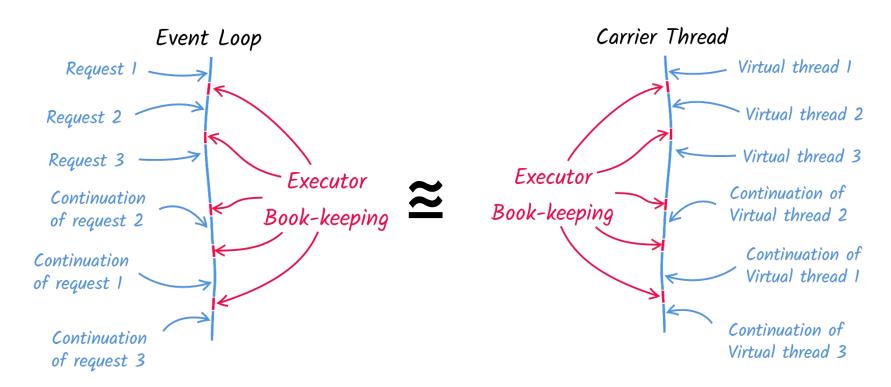
ops/s

(executionStrategy)	Mode	Cnt
Native	thrpt	40
ForkJoinVirtual	thrpt	40
FixedCarrierPoolVirtual	thrpt	40
PinnedCarrierVirtual	thrpt	40

So why we may want to plug a different carrier? The Quarkus use case - Reactive ≅ Virtual Threads

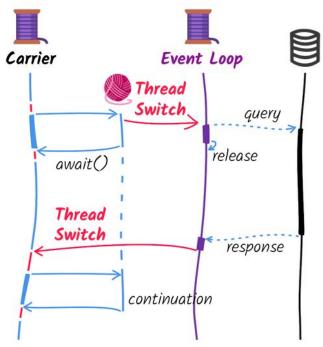


So why we may want to plug a different carrier? The Quarkus use case - Reactive ≅ Virtual Threads

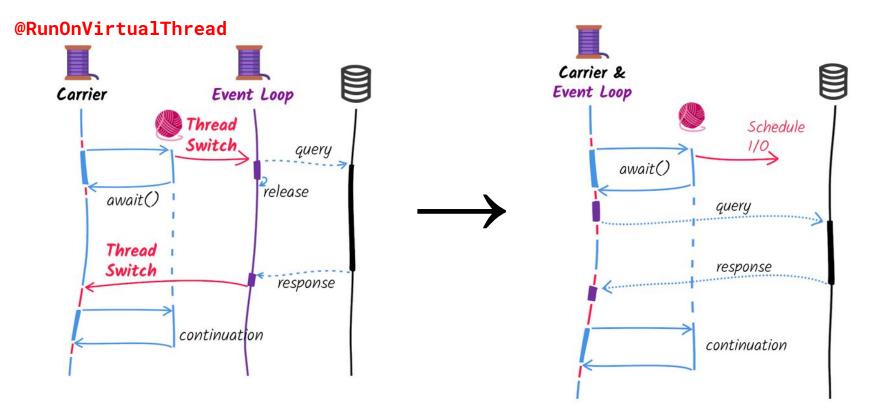


So why we may want to plug a different carrier? The Quarkus use case - The Event Loop as carrier

@RunOnVirtualThread



So why we may want to plug a different carrier? The Quarkus use case - The Event Loop as carrier



Why not using the event-loop as carrier thread???

No (more) API to do that in Loom - API removed in fall 2021

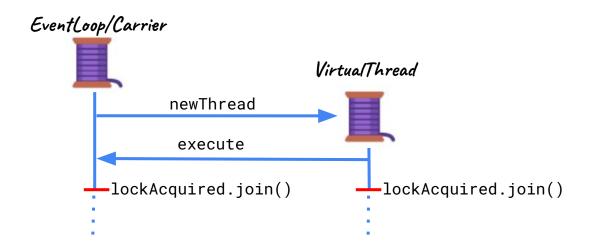
Why not using the event-loop as carrier thread???

- No (more) API to do that in Loom API removed in fall 2021
- It lead to.... deadlocks!

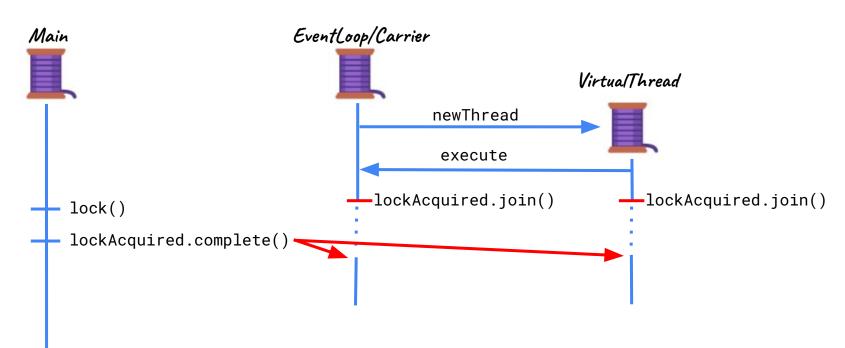




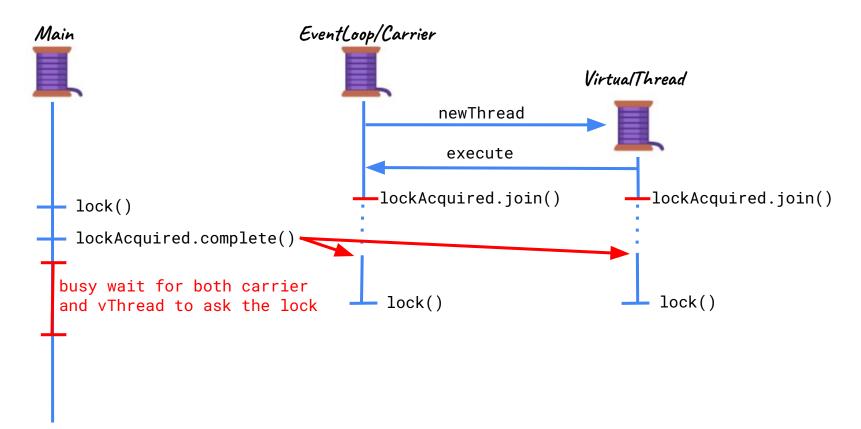




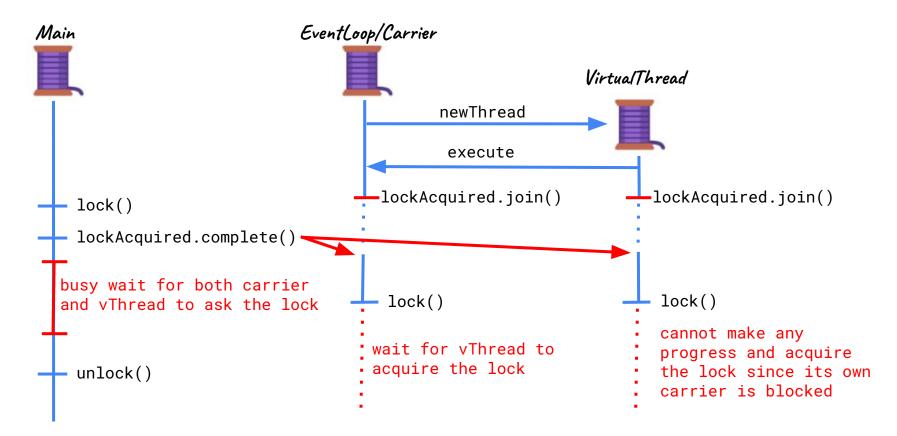












Conclusions

- ❖ Virtual threads are **not faster** threads ⇒ they don't magically execute more instructions per second than native ones do.
- ❖ Virtual threads are not drop in replacement for native threads ⇒ they are a different (new) tool and you have to know their features and characteristics in order to use them appropriately.
- ❖ Virtual threads are generally not a good fit for CPU bound tasks ⇒ lack of fairness and the performance cost of coroutines can be problematic.
- ♦ What virtual threads are really good for is waiting ⇒ their goal is maximizing the utilization of external resources and then improving throughput, without affecting readability.

Learn More

Quarkus Virtual Threads

- REST / HTTP
- Hibernate, Bean Validation, Transactions
- Kafka, AMQP, JMS
- gRPC
- Scheduled tasks
- Event Bus
- ...

